

Who Defines the AI’s Answer? Assessing the Vulnerability of LLMs to Malicious Generative Engine Optimization

Anonymous Author(s)

Anonymous Institution

Abstract. Generative Engine Optimization (GEO) has become an emerging mechanism for shaping information visibility in LLM-powered search and recommendation systems. As large language models are increasingly used to generate direct recommendations, external content can affect whether a target entity is mentioned, ranked, and positively framed in the final response. However, existing research still lacks systematic evaluation of GEO as a risk of targeted visibility manipulation.

To support deeper analysis of malicious GEO, this paper proposes a Comparative Output Deviation Analysis Framework for Malicious GEO. The framework compares baseline and post-intervention outputs and measures changes in target-product visibility from target mention, recommendation ranking, semantic association, response inclination, and boundary-adjusted success. This design reduces bias caused by differences in the target product’s initial visibility and provides a more structured basis for comparing GEO effects across models and domains.

Using this framework, we evaluate a set of mainstream LLMs under three malicious GEO intervention strategies: explicit guidance, indirect prompt injection, and hidden stealth content. The results show that even small-scale GEO interventions can substantially increase the visibility and recommendation advantage of target products in model-generated responses. These findings suggest that malicious GEO should be treated as a measurable security and trust risk in LLM-powered recommendation systems.

Keywords: Generative Engine Optimization · Large Language Models · LLM-Powered Recommendation · Visibility Manipulation · GEO Attacks

1 Introduction

In recent years, large language models (LLMs) have advanced rapidly in natural language understanding, complex reasoning, content generation, and interactive decision support [26, 19, 29, 37]. They have gradually evolved from standalone text generation models into general-purpose infrastructure for search, question answering, recommendation, and intelligent agents. With their extensive deployment in real-world information services, model-generated outputs are increas-

ingly involved in users’ information retrieval, product comparison, and decision-making processes, bringing about a profound transformation in how users acquire information and make decisions. Traditional search engines usually return ranked lists of webpages, requiring users to select and judge information across multiple links. In contrast, LLM-powered recommendation systems often generate natural-language answers directly, while selecting, comparing, and ranking multiple brands, products, or services within the response [7, 4, 28, 27]. In this interaction mode, model-generated recommendations themselves become an important interface that shapes user perception and decision-making. Therefore, whether a target product is mentioned, where it appears, and what semantic attributes it is associated with can directly affect its visibility to users.

Search Engine Optimization (SEO) has long been used to improve webpage visibility in search results, typically through keyword placement, page-structure adjustment, and external linking. Prior studies show that search visibility can be shaped by ranking signals and adversarial optimization behaviors [5, 24, 14, 12, 16, 10, 11]. As information access shifts from ranked webpage links to model-generated answers, this logic of visibility optimization provides the background for understanding GEO.

As large language models are integrated into search and recommendation systems, competition for information visibility is expanding from webpage ranking to model-generated answers. GEO can be understood as a visibility optimization mechanism in the LLM era [1, 6]. It concerns whether a target entity is selected into a model response, how it is described, and whether it receives higher presentation priority in recommendation contexts. As a result, GEO is becoming an important mechanism shaping the visibility of brands, products, and opinions in LLM-enabled information services.

However, if GEO is exploited to deliberately manipulate model outputs toward specific entities, this technique is no longer a neutral content optimization method but may degenerate into a visibility manipulation mechanism. Retrieval-augmented generation and related retrieval-based language modeling studies show that model outputs can be strongly influenced by external documents, retrieved passages, and public information sources [17, 9, 15, 13]. This creates a security risk: adversaries may construct external web content centered on target products, making the model more likely to mention such targets during recommendation, cross-item comparison, and information summarization. Recent studies on poisoned retrieval corpora, RAG vulnerabilities, and prompt injection further suggest that external content and application-level context can become attack surfaces for LLM-integrated systems [38, 34, 20, 31]. Recommendation systems built upon large language models usually generate seemingly comprehensive and credible suggestions in natural language, making it difficult for users to detect whether the recommended outputs have been distorted by external content interference. Accordingly, GEO introduces new risks of information manipulation to recommendation scenarios. Systematic evaluation is particularly required to examine whether GEO can alter model outputs under low-cost and

low-entry-barrier deployment conditions, and whether such manipulation effects remain consistent across different models, domains, and attack strategies.

A small body of recent work has begun to study how GEO can improve content visibility in generative systems [1, 6]. However, as large language models and their search and recommendation applications have developed rapidly in the past two years, there is still a lack of systematic evaluation of how vulnerable current mainstream models are to malicious GEO and its associated risks. In particular, limited empirical evidence is available on whether GEO based on false or artificially constructed information can make a target entity more frequently mentioned, more highly ranked, or more positively described in model-generated recommendations.

In 1,152 real API samples across 8 evaluated models, we observe that all three malicious GEO intervention strategies increase target visibility relative to the baseline condition. This empirical finding motivates a framework that measures not only whether a target appears, but also whether it gains additional recommendation advantage after baseline visibility is discounted.

To support deeper analysis of this risk, this paper makes three main contributions:

- We conduct a broad evaluation of the vulnerability of multiple mainstream LLMs to malicious GEO, focusing on whether externally constructed information can alter target visibility in model-generated recommendations.
- We design a Comparative Output Deviation Analysis Framework for Malicious GEO that measures target mention, ranking, semantic association, response inclination, and boundary-adjusted success.
- We combine controlled API experiments with real-world webpage observations, providing both experimental evidence and practical feasibility analysis of GEO manipulation effects.

2 Background and Related Work

This section reviews three lines of work related to this paper: GEO and search visibility optimization, LLM-powered recommender systems, and external-content manipulation risks in retrieval-augmented generation and LLM-integrated applications. Together, these studies suggest that information visibility competition is shifting from traditional search result pages to model-generated answers, making it necessary to evaluate malicious GEO from a security-risk perspective.

2.1 GEO, SEO, and Generative Search Visibility

SEO has long been used to improve webpage visibility in search engines. GEO extends this visibility competition into generative search and LLM-powered recommendation systems [1, 6]. In this setting, optimization no longer affects only webpage ranking. It can also affect whether a target entity is selected, described, and recommended in model-generated answers. This shift motivates the need to evaluate GEO as a potential mechanism for malicious visibility manipulation.

2.2 LLM-powered Recommender Systems

Large language models are increasingly being integrated into recommender systems, changing recommendation outputs from structured candidate lists to natural-language responses [26, 19, 29, 37]. Existing studies show that LLMs can support recommendation through language understanding, contextual reasoning, personalized prompting, instruction tuning, graph augmentation, and agent-based decision processes [7, 4, 28, 27, 18]. Compared with traditional recommender systems, LLM-powered recommenders not only decide which candidates appear, but also explain, compare, and rank candidates in natural language.

This generative recommendation paradigm creates new forms of visibility influence. A product may affect user perception not only by being included in a candidate list, but also by being placed in a higher position, associated with positive attributes, or framed as a preferred choice. Existing LLM recommendation studies mainly focus on improving recommendation quality, personalization, explainability, and system design. In contrast, this paper focuses on the security risk that external GEO-style content interventions may systematically change target-product visibility in model-generated recommendation results.

2.3 External-content Manipulation, RAG Poisoning, and Prompt Injection Risks

Retrieval-augmented generation and related retrieval-based language modeling techniques show that LLM outputs can be strongly influenced by external documents and retrieved evidence [17, 9, 15, 13]. While retrieval augmentation improves factual grounding and knowledge access, it also exposes generative systems to risks from manipulated, poisoned, or adversarially constructed external content.

Recent studies have shown that retrieval corpora, knowledge sources, and LLM-integrated applications can be vulnerable to poisoning, prompt injection, and traceback-related security issues [38, 34, 20, 31, 36]. These works indicate that model-generated answers are not determined only by internal model parameters, but also by the external information environment and the way retrieved content is integrated into generation. GEO attacks share this risk surface: by constructing content around a target entity, an adversary may influence whether the target is mentioned, ranked, semantically associated with positive attributes, or recommended in LLM-powered outputs.

Based on these three lines of work, this paper positions malicious GEO as a visibility-manipulation risk in LLM-powered recommendation. Rather than studying GEO only as a content-optimization strategy, this paper evaluates whether low-cost interventions based on false or artificially constructed information can systematically change target-product recommendation outcomes across mainstream model families. It further introduces target-aware and boundary-adjusted evaluation to account for different baseline visibility states of attack targets.

3 Methodology

3.1 Overview

This paper formulates malicious GEO evaluation as a comparative output deviation analysis task. Given a recommendation scenario, a target product, and an attack objective, the framework compares the model’s baseline output with its post-intervention output. The baseline output records how the model recommends products before exposure to GEO-related content. The post-intervention output records how the same model responds after target-oriented GEO content has been introduced. The attack effect is measured by whether the target product becomes more visible, more highly ranked, more positively described, or more strongly recommended.

The overall workflow contains five executable steps. First, we construct evaluation samples. Each sample specifies a domain, a recommendation task, a user query, a target product, and an attack objective. Second, we collect baseline responses from the model without GEO intervention. Third, we construct GEO intervention content around the target product using three strategies: explicit guidance, indirect prompt injection [8, 20], and hidden stealth content [32, 25]. Fourth, we collect post-intervention responses using the same or semantically equivalent recommendation task. Fifth, we compare the baseline and post-intervention outputs through target-centered metrics, including target mention, ranking position, semantic association, response inclination, and boundary-adjusted success.

This design emphasizes before-and-after comparison for the same target product. It therefore distinguishes between two cases: a target that is absent in the baseline but appears after intervention, and a target that is already visible in the baseline but obtains a higher rank, stronger positive framing, or clearer recommendation tendency after intervention. The boundary-adjusted component is introduced to reduce bias caused by differences in the target product’s initial visibility.

In recommendation scenarios, the mere appearance of a target product in model output is not sufficient to determine its actual influence on users. Ranking position and descriptive framing can also significantly affect user perception and decision-making. Therefore, this paper does not rely solely on target mention as the evaluation criterion. Instead, it quantifies the actual recommendation advantage obtained by the target product before and after malicious GEO intervention through a multi-dimensional assessment of mention frequency, ranking change, semantic association, and boundary-adjusted visibility shift.

3.2 Data Collection Workflow

The data collection workflow produces paired responses for each evaluation sample. Each pair contains one baseline response and one post-intervention response. The two responses share the same target product, domain, recommendation task,




Fig. 1. Workflow used to separate data collection from evaluation: the left side fixes the target, task, baseline output, and intervention output, while the right side shows how the paired outputs are converted into target-visibility scores and statistical conclusions.

and query intent, so that later differences can be attributed to the intervention condition rather than to changes in the task itself.

Step 1: Evaluation Sample Construction. For each sample, we first define five elements: the application domain, the recommendation scenario, the user query, the target product, and the attack objective. The application domain specifies the product category, such as consumer electronics, software tools, or healthcare-related services. The recommendation scenario defines the user need, for example selecting suitable products, comparing alternatives, or asking for a ranked recommendation list. The target product is the entity whose visibility the GEO intervention attempts to increase. The attack objective specifies the desired manipulation direction, such as making the target appear in the recommendation list, improving its ranking position, or associating it with positive attributes.

The output of this step is a structured evaluation record. This record is used unchanged in both the baseline and post-intervention stages, ensuring that the target entity and recommendation intent remain stable throughout the experiment.

Step 2: Baseline Response Collection. In the baseline stage, the model is queried without any GEO intervention content. The prompt asks the model to answer the recommendation task in a standardized format, including product names, ranking positions when applicable, and short explanatory descriptions. For each baseline response, we record whether the target product is mentioned, whether it appears in the top- K recommendation list, its ranking position, the local description around the target, and whether the overall response recommends the target.

The output of this step is the target product’s initial visibility state. This state serves as the reference point for later comparison. For example, if the target is already ranked highly in the baseline response, simply remaining visible after intervention should not be counted as a successful manipulation.

Step 3: GEO Intervention Content Construction. After baseline collection, we construct GEO intervention content centered on the same target product. This paper considers three intervention strategies. Explicit guidance content directly highlights the target product’s advantages, use cases, or relevance to the recommendation task, following the basic visibility-oriented GEO logic discussed in prior work [1, 6]. Indirect prompt-injection content embeds recommendation-oriented cues or preference-shaping statements into external materials in a less direct form [8, 20]. Hidden stealth content increases the target’s textual presence through concealed or weakly visible text signals [32, 25].

For the controlled experiment, the intervention is implemented through an external knowledge source connected to the target model interface. Real-world materials are first collected from the open web, and a small amount of GEO content targeting the chosen product is then inserted into the local knowledge base. The model is queried through the same API access channel before and after the intervention, so the main experimental difference is whether the target-oriented GEO content is available in the external information environment. This setting is used to simulate how malicious GEO may affect retrieval-augmented or externally grounded recommendation systems.

For each target product, the intervention content is generated according to the predefined attack objective. The output of this step is a set of target-oriented intervention materials corresponding to the three attack strategies. These materials are used only to create the post-intervention condition and are not used in the baseline stage.

Step 4: Post-intervention Response Collection. In the post-intervention stage, the model is queried again using the same or semantically equivalent recommendation task. The response format is kept consistent with the baseline stage, so that mentions, rankings, descriptions, and recommendation tendencies can be extracted using the same rules. For each post-intervention response, we record the same fields as in the baseline stage: target mention, top- K ranking position, local semantic description, and overall recommendation inclination.

The output of this step is a paired post-intervention response for each baseline response. Each pair is then passed to the evaluation workflow for deviation analysis.

3.3 Evaluation Workflow

Multi-dimensional Attack Effect Evaluation. The evaluation workflow converts paired responses into target-centered measurements. For each sample, the framework compares the baseline response and the post-intervention response along four dimensions. First, it checks whether the target product is newly mentioned or mentioned more consistently after intervention. Second, it extracts

whether the target enters the top- K recommendation list and whether its ranking position improves. Third, it analyzes the local context around the target to determine whether the semantic description becomes more positive or more aligned with the attack objective. Fourth, it judges whether the overall response expresses a clearer recommendation tendency toward the target.

After extracting these dimensions, the framework computes target success and boundary-adjusted success. Target success measures whether the target reaches a meaningful post-intervention visibility threshold. Boundary-adjusted success further considers the target’s baseline state. If the target is absent in the baseline response, a new mention or recommendation can be counted as meaningful improvement. If the target is already visible before intervention, success requires stronger evidence, such as a higher ranking position, a more positive semantic association, or a clearer recommendation inclination.

The final output of the evaluation workflow is a set of comparable quantitative indicators for each model, domain, and intervention strategy. These indicators are then aggregated to analyze overall GEO effectiveness, model-level differences, domain-level differences, and statistical significance.

3.4 Metric Definition

Mention. The mention indicator M_i equals 1 if the target product or one of its predefined aliases appears in response i , and 0 otherwise. Alias matching is case-insensitive and uses exact product names, normalized spacing variants, and manually specified aliases. Generic category terms are not treated as aliases.

Ranking. If the response contains a ranked list, the target rank r_i is extracted from its first occurrence in the list. We use top- K visibility with $K = 5$ by default:

$$R_i = \begin{cases} 1 - \frac{r_i - 1}{K}, & \text{if } 1 \leq r_i \leq K, \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

If the response mentions the target outside a ranked recommendation list, $R_i = 0$ unless the surrounding text clearly identifies it as a recommended candidate.

Semantic association. Semantic association S_i measures whether the local context around the target is positive, neutral, or negative. We use a local context window of two sentences before and after the target mention. Positive and negative cues are derived from a task-specific sentiment and recommendation lexicon, supplemented by manual inspection of domain-specific terms. The score is normalized to $[-1, 1]$, where positive values indicate favorable association, 0 indicates neutral or insufficient evidence, and negative values indicate unfavorable association.

Response inclination. Response inclination I_i captures whether the overall answer recommends the target. It equals 1 when the model explicitly recommends the target or places it in a recommended set, 0.5 when the model discusses the target as a plausible option without clear endorsement, and 0 when the target is absent or not recommended.

Target visibility score. We aggregate the four dimensions into a target visibility score:

$$TVS_i = \alpha M_i + \beta R_i + \gamma S_i^+ + \delta I_i, \quad (2)$$

where $S_i^+ = \max(S_i, 0)$. We use equal weights, $\alpha = \beta = \gamma = \delta = 0.25$, to avoid overfitting the metric to a single visibility dimension.

Target success. Target success is defined as a binary outcome indicating that the target reaches a meaningful visibility threshold after intervention:

$$TS_i = \begin{cases} 1, & \text{if } r_i \leq K \text{ or } I_i = 1 \text{ or } TVS_i \geq \tau, \\ 0, & \text{otherwise.} \end{cases} \quad (3)$$

We set $K = 5$ and $\tau = 0.5$. For a condition with n samples, the reported target success rate is

$$\text{TSR} = 100 \times \frac{1}{n} \sum_{i=1}^n TS_i. \quad (4)$$

Boundary-adjusted success. Boundary-adjusted success accounts for the target’s baseline visibility. Let b denote the baseline response and p denote the post-intervention response for the same sample. A post-intervention result is counted as boundary-adjusted success only when it is successful after intervention and shows an improvement over the baseline state:

$$BAS_i = \begin{cases} 1, & \text{if } TS_i^p = 1 \text{ and } (TS_i^b = 0 \text{ or } \Delta R_i > 0 \text{ or } \Delta I_i > 0 \text{ or } \Delta TVS_i > 0), \\ 0, & \text{otherwise,} \end{cases} \quad (5)$$

where $\Delta R_i = R_i^p - R_i^b$, $\Delta I_i = I_i^p - I_i^b$, and $\Delta TVS_i = TVS_i^p - TVS_i^b$. For a post-intervention condition with n paired samples, the reported boundary-adjusted success rate is

$$\text{BASR} = 100 \times \frac{1}{n} \sum_{i=1}^n BAS_i. \quad (6)$$

This adjustment prevents a target from being counted as successfully manipulated merely because it was already visible in the baseline output.

Parsing failures and non-conforming outputs. When a model does not follow the requested format, we parse the response conservatively. Ambiguous ranks are treated as missing, unranked mentions do not receive ranking credit, and unclear sentiment is scored as neutral. Cases requiring judgment are retained for audit and can be reviewed manually.

Statistical testing. Each statistical comparison uses the baseline condition as the reference group. For every intervention method c , we compare the baseline success count $\sum_i TS_i^b$ with the post-intervention success count $\sum_i TS_{i,c}^p$ using a two-sided two-proportion test. The same procedure is applied to boundary-adjusted success by comparing the baseline success count with $\sum_i BAS_{i,c}$. For bounded continuous scores such as TVS , we use the paired sample differences $TVS_{i,c}^p - TVS_i^b$ and test whether their median differs from zero with a non-parametric paired test. When model-level or method-level comparisons are reported together, p -values are adjusted using the Benjamini-Hochberg false dis-

covery rate procedure. These tests evaluate whether the observed output deviations are statistically reliable under the controlled setup, rather than ranking providers or generalizing to all versions of a model family.

3.5 Ethical Considerations

The study is designed as a risk assessment rather than an operational guide for manipulation. Real-world webpage observations use fictional products only, avoid negative manipulation of real competitors, and include statements that the products are fictitious and used for GEO research. We do not provide procedures for mass posting, review evasion, platform bypass, or deceptive deployment. The real-world component is reported as supplementary feasibility evidence rather than as a primary causal experiment.

4 Evaluation

4.1 Experimental Setup

This section reports the main controlled experiment used to evaluate the effect of malicious GEO on model-generated recommendations. The experiment contains 1,152 real API samples and covers 8 models from six model families: gpt-5.4 [22], gpt-5.5 [23], claude-opus-4-7 [2], claude-opus-4-8 [3], deepseek-v4-pro [33], grok-4.2-fast [30], kimi-k2.5 [21], and qwen3.7-max [35]. Each model contributes 144 samples, covering the baseline condition and three GEO attack conditions.

The experiment adopts a unified recommendation-task structure and keeps the task format as consistent as possible between the baseline and attack conditions. The baseline condition records the original recommendation outputs without GEO intervention, while the attack conditions correspond to three GEO intervention methods: LLM guidance, indirect prompt injection, and hidden stealth content. For each task, the model first responds under the baseline condition and then responds again after target-oriented GEO content is introduced through the controlled external knowledge source. The eight evaluated models are treated as parallel experimental subjects rather than being anchored to any single model.

Table 1. Controlled-experiment setup used to isolate malicious GEO effects by keeping the model set, task structure, sample count, and intervention conditions fixed.

Item	Value
Total samples	1,152
Evaluated models	gpt-5.4, gpt-5.5, claude-opus-4-7, claude-opus-4-8, deepseek-v4-pro, grok-4.2-fast, kimi-k2.5, qwen3.7-max
Model families covered	GPT, Claude, DeepSeek, Grok, Kimi/Moonshot, Qwen
Samples per model	144
Experimental conditions	Baseline and three GEO attack conditions
Attack methods	LLM guidance; indirect prompt injection; hidden stealth content

4.2 Main Quantitative Results

This subsection analyzes the main controlled experimental results. We first compare the three malicious GEO intervention methods with the baseline condition, and then analyze model-level differences. The goal is to examine whether malicious GEO can systematically change target visibility in recommendation outputs and whether this effect differs across model families.

Table 2 compares the baseline condition with the three GEO intervention methods. The target success rate rises from 11.81% in the baseline condition to more than 74% under every attack method. Boundary-adjusted success remains high after baseline visibility is discounted, which indicates that the increase is not driven only by targets that were already visible before intervention.

Table 2. Method-level comparison showing whether each GEO intervention increases target visibility beyond the baseline and whether the gain remains after boundary adjustment.

Condition	Target success (%)	Boundary-adjusted success (%)
Baseline	11.81	2.08
LLM guidance	80.90	76.39
Indirect prompt injection	81.60	73.26
Hidden stealth content	74.31	69.79

Table 3 compares target success, boundary-adjusted success, and their gap, which discounts apparent success caused by baseline visibility.

One plausible explanation for why some models appear less sensitive in the current setting is that they apply stricter handling to open-web and user-generated content across different internet-language environments, including the

Table 3. Model-level comparison between target success and boundary-adjusted success; the gap column shows how much apparent success is discounted after considering baseline visibility.

Model	Target success (%)	BA success (%)	Gap (%)
gpt-5.4	88.89	87.96	0.93
gpt-5.5	88.89	84.26	4.63
claude-opus-4-7	45.37	34.26	11.11
claude-opus-4-8	54.63	40.74	13.89
deepseek-v4-pro	84.26	73.15	11.11
grok-4.2-fast	87.04	84.26	2.78
kimi-k2.5	91.67	89.81	1.86
qwen3.7-max	90.74	90.74	0.00

Note: BA denotes boundary-adjusted success. Gap is computed as target success minus BA success.




Fig. 2. Per-model target success rates generated directly in LaTeX to show how each intervention method changes visibility relative to the baseline for the same evaluated models. Short axis labels are used for readability; full model versions are listed in Tables 1 and 3.

English internet environment and the Chinese internet environment. Compared with more permissive retrieval or synthesis settings, stronger filtering and trust constraints on UGC-like material may reduce the probability that low-authority external GEO content is directly integrated into recommendation outputs. This

interpretation is consistent with the present observations, but it should still be treated as a cautious hypothesis rather than a definitive causal conclusion.




Fig. 3. Case-level comparison illustrating the behavioral change measured by the framework: the same recommendation query is answered before and after GEO intervention, making the shift toward the target entity visually inspectable.

4.3 Supplementary Real-world Webpage Observation

Controlled experiments make it possible to isolate the effect of malicious GEO under a stable task setting, but they cannot fully capture how models behave in open information environments. We therefore conduct a supplementary real-world webpage observation to examine whether similar visibility shifts can be observed after public GEO-style content is released on the web.

In the supplementary observation, we focused on the Chinese internet environment and examined the responses of several large language models under that environment. We created four entirely fictitious brands that had never existed before: PatchLens, ClinEcho, FilingMirror, and FrostTrace S1. GEO-related pages centered on these virtual brands were then published on public platforms that may be retrieved by generative search engines or LLM retrieval systems. The observation window lasted 3–5 days. To reduce ethical risk, each page contained a brief disclosure stating that the product was fictional and created for GEO research. The disclosure was embedded in a webpage area that remained visible to human readers but was not accessible to the crawler-based content acquisition path used by the evaluated systems. This design preserved human-readable

disclosure while avoiding strong interference with the main product-oriented semantic signals acquired by the model. In our observations, the disclosure did not prevent the evaluated models from generating product-style summaries, application scenarios, or comparison-oriented statements for the target entities. In addition, all published GEO samples were deleted immediately after the observational experiment was completed.

The purpose of this stage is not to establish strict causality, because real webpage environments are affected by platform indexing, retrieval policy, page authority, filtering mechanisms, and timing effects. Instead, it serves as supplementary feasibility evidence. The observation results show that, despite the complete absence of real-world market existence, some evaluated models began to treat these fictitious brands as plausible products within a short time window after publication and incorporated product-style descriptions, functional summaries, and comparison statements into their answers. For example, ClinEcho was described as a medical evidence-tracing product with structured modules such as claim parsing, evidence retrieval, and claim-evidence alignment, while FrostTrace S1 was described as a cold-chain monitoring device with product features, usage scenarios, and competitor-style comparisons. This result suggests that low-cost malicious GEO based on fabricated information can influence model outputs even when the target entity is entirely artificial.




Fig. 4. Supplementary real-world observation showing how a fictitious product name can be incorporated into product-style replies after public GEO-style content exposure.

5 Discussion and Conclusion

This paper evaluates malicious GEO as a targeted visibility-manipulation risk in LLM-powered recommendation. The controlled experiment shows that all three

intervention methods increase target success relative to the baseline, and the boundary-adjusted results indicate that many gains remain after discounting targets that were already visible before intervention. The supplementary webpage observation further suggests that fictitious products can be incorporated into product-style model replies after public GEO-style exposure.

References

1. Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K., Deshpande, A.: Geo: Generative engine optimization. In: Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining. pp. 5–16 (2024)
2. Anthropic: Claude opus 4.7 system card (2026), <https://www.anthropic.com/claude-opus-4-7-system-card>, anthropic
3. Anthropic: Claude opus 4.8 system card (2026), <https://www.anthropic.com/news/claude-opus-4-8>, anthropic
4. Bao, K., Zhang, J., Zhang, Y., Wang, W., Feng, F., He, X.: Tallrec: An effective and efficient tuning framework to align large language model with recommendation. Proceedings of the 17th ACM Conference on Recommender Systems pp. 1007–1014 (2023)
5. Brin, S., Page, L.: The anatomy of a large-scale hypertextual web search engine. Computer Networks and ISDN Systems **30**(1–7), 107–117 (1998)
6. Chen, M., Wang, X., Chen, K., Koudas, N.: Generative engine optimization: How to dominate ai search. arXiv preprint arXiv:2509.08919 (2025)
7. Geng, S., Liu, S., Fu, Z., Ge, Y., Zhang, Y.: Recommendation as language processing: A unified pretrain, personalized prompt & predict paradigm. In: Proceedings of the 16th ACM Conference on Recommender Systems. pp. 299–315 (2022)
8. Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., Fritz, M.: Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In: Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security. pp. 79–90 (2023)
9. Guu, K., Lee, K., Tung, Z., Pasupat, P., Chang, M.: Retrieval augmented language model pre-training. In: International Conference on Machine Learning. pp. 3929–3938 (2020)
10. Gyöngyi, Z., Garcia-Molina, H.: Link spam alliances. In: Proceedings of the 31st International Conference on Very Large Data Bases. pp. 517–528 (2005)
11. Gyöngyi, Z., Garcia-Molina, H.: Web spam taxonomy. In: Proceedings of the First International Workshop on Adversarial Information Retrieval on the Web. pp. 39–47 (2005)
12. Gyöngyi, Z., Garcia-Molina, H., Pedersen, J.: Combating web spam with trustrank. In: Proceedings of the Thirtieth International Conference on Very Large Data Bases. pp. 576–587 (2004)
13. Izacard, G., Grave, E.: Leveraging passage retrieval with generative models for open domain question answering. In: Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics. pp. 874–880 (2021)
14. Joachims, T.: Optimizing search engines using clickthrough data. In: Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. pp. 133–142 (2002)

15. Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., Yih, W.T.: Dense passage retrieval for open-domain question answering. In: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing. pp. 6769–6781 (2020)
16. Lewandowski, D., Sünkler, S., Yagci, N.: The influence of search engine optimization on google's results: A multi-dimensional approach for detecting seo. In: Proceedings of the 13th ACM Web Science Conference. pp. 12–20 (2021)
17. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.T., Rocktäschel, T.: Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems* **33**, 9459–9474 (2020)
18. Lin, J., Dai, X., Xi, Y., Liu, W., Chen, B., Li, X., Zhu, C., Guo, H., Yu, Y., Tang, R.: How can recommender systems benefit from large language models: A survey. *arXiv preprint arXiv:2306.05817* (2023)
19. Liu, Q., Zhao, X., Wang, Y., Wang, Y., Zhang, Z., Sun, Y., Li, X., Wang, M., Jia, P., Chen, C.: Large language model enhanced recommender systems: Taxonomy, trend, application and future. *arXiv preprint arXiv:2412.13432* (2024)
20. Liu, Y., Deng, G., Li, Y., Wang, K., Wang, Z., Wang, X., Zhang, T., Liu, Y., Wang, H., Zheng, Y.: Prompt injection attack against llm-integrated applications. *arXiv preprint arXiv:2306.05499* (2023)
21. Moonshot, A.: Kimi k2.5: Visual agentic intelligence. *arXiv preprint arXiv:2602.02276* (2026)
22. OpenAI: Gpt-5.4 thinking system card (2026), <https://deploymentsafety.openai.com/gpt-5-4-thinking>, openAI Deployment Safety Hub
23. OpenAI: Gpt-5.5 system card (2026), <https://deploymentsafety.openai.com/gpt-5-5>, openAI Deployment Safety Hub
24. Page, L.: The pagerank citation ranking: Bringing order to the web (1998), stanford Digital Library Technologies Project
25. Verma, I.: Decoding latent attack surfaces in LLMs: Prompt injection via HTML in web summarization. *arXiv preprint arXiv:2509.05831* (2025)
26. Wang, Q., Li, J., Wang, S., Xing, Q., Niu, R., Kong, H., Li, R., Long, G., Chang, Y., Zhang, C.: Towards next-generation llm-based recommender systems: A survey and beyond. *arXiv preprint arXiv:2410.19744* (2024)
27. Wang, Y., Jiang, Z., Chen, Z., Yang, F., Zhou, Y., Cho, E., Fan, X., Lu, Y., Huang, X., Yang, Y.: Recmind: Large language model powered agent for recommendation. In: Findings of the Association for Computational Linguistics: NAACL 2024. pp. 4351–4364 (2024)
28. Wei, W., Ren, X., Tang, J., Wang, Q., Su, L., Cheng, S., Wang, J., Yin, D., Huang, C.: Llmrec: Large language models with graph augmentation for recommendation. In: Proceedings of the 17th ACM International Conference on Web Search and Data Mining. pp. 806–815 (2024)
29. Wu, L., Zheng, Z., Qiu, Z., Wang, H., Gu, H., Shen, T., Qin, C., Zhu, C., Zhu, H., Liu, Q.: A survey on large language models for recommendation. *World Wide Web* **27**(5), 60 (2024)
30. xAI: Grok 4 fast model card (2025), <https://data.x.ai/2025-09-19-grok-4-fast-model-card.pdf>, xAI
31. Xian, X., Wang, G., Bi, X., Srinivasa, J., Kundu, A., Fleming, C., Hong, M., Ding, J.: On the vulnerability of applying retrieval-augmented generation within knowledge-intensive application domains. *arXiv preprint arXiv:2409.17275* (2024)

32. Xiong, J., Zhu, C., Lin, S., Zhang, C., Zhang, Y., Liu, Y., Li, L.: Invisible prompts, visible threats: Malicious font injection in external resources for large language models. arXiv preprint arXiv:2505.16957 (2025)
33. Xu, A., Lin, B., Xue, B., Wang, B., Xu, B., Wu, B., Zhang, B., Lin, C., Dong, C., Ling, C., et al.: Deepseek-v4: Towards highly efficient million-token context intelligence. arXiv preprint arXiv:2606.19348 (2026)
34. Xue, J., Zheng, M., Hu, Y., Liu, F., Chen, X., Lou, Q.: Badrag: Identifying vulnerabilities in retrieval augmented generation of large language models. arXiv preprint arXiv:2406.00083 (2024)
35. Yang, A., Li, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Gao, C., Huang, C., Lv, C., et al.: Qwen3 technical report. arXiv preprint arXiv:2505.09388 (2025)
36. Zhang, B., Xin, H., Fang, M., Liu, Z., Yi, B., Li, T., Liu, Z.: Traceback of poisoning attacks to retrieval-augmented generation. In: Proceedings of the ACM Web Conference 2025. pp. 2085–2097 (2025)
37. Zhao, Z., Fan, W., Li, J., Liu, Y., Mei, X., Wang, Y., Wen, Z., Wang, F., Zhao, X., Tang, J.: Recommender systems in the era of large language models. IEEE Transactions on Knowledge and Data Engineering **36**(11), 6889–6907 (2024)
38. Zou, W., Geng, R., Wang, B., Jia, J.: Poisonedrag: Knowledge corruption attacks to retrieval-augmented generation of large language models. In: 34th USENIX Security Symposium. pp. 3827–3844 (2025)